

LightGBM Classifier in Python

LightGBM is a fast, distributed, high performance gradient boosting framework based on decision tree algorithms, used for ranking, classification and many other machine learning tasks.

Table of Contents

- [1. Introduction to LightGBM](#)
- [2. LightGBM intuition](#)
- [3. XGBoost Vs LightGBM](#)
- [4. LightGBM Parameters](#)
 - [4.1 Control Parameters](#)
 - [4.2 Core Parameters](#)
 - [4.3 Metric Parameter](#)
 - [4.4 IO Parameter](#)
- [5. LightGBM implementation in Python](#)
- [6. LightGBM Parameter Tuning](#)
- [7. References](#)

1. Introduction to LightGBM

- [LightGBM](#) is a gradient boosting framework that uses tree based learning algorithms
- Faster training speed and higher efficiency.
 - Lower memory usage.
 - Better accuracy.
 - Support of parallel and GPU learning.
 - Capable of handling large-scale data.
- At present, decision tree based machine learning algorithms dominate Kaggle competitions. The winning solutions in these competitions have adopted an algorithm called **XGBoost**.

Light GBM can handle the large size of data and takes lower memory to run.

- Another reason why Light GBM is so popular is because it focuses on accuracy of results. LGBM also supports GPU learning and thus data scientists are widely using LGBM for data science application development.
- It is not advisable to use LGBM on small datasets. Light GBM is sensitive to overfitting and can easily overfit small data.

3. XGBoost Vs LightGBM

Table of Contents

- [XGBoost](#) is a very fast and accurate ML algorithm. But now it's been challenged by [LightGBM](#) — which runs even faster with comparable model accuracy and more hyperparameters for users to tune.
- The key difference in speed is because **XGBoost split the tree nodes one level at a time** and **LightGBM does that one node at a time**.

4. LightGBM Parameters

4.1 Control Parameters

4.2 Core Parameters

4.3 Metric Parameter

4.4 IO Parameter

4.1 Control Parameters

Table of Contents

- **max_depth** : It describes the maximum depth of tree. This parameter is used to handle model overfitting. If you feel that your model is overfitted, you should to lower max_depth.
- **min_data_in_leaf** : It is the minimum number of the records a leaf may have. The default value is 20, optimum value. It is also used to deal with overfitting.
- **feature_fraction**: Used when your boosting is random forest. 0.8 feature fraction means LightGBM will select 80% of parameters randomly in each iteration for building trees.
- **bagging_fraction** : specifies the fraction of data to be used for each iteration and is generally used to speed up the training and avoid overfitting.
- **early_stopping_round** : This parameter can help you speed up your analysis. Model will stop training if one metric of one validation data doesn't improve in last early_stopping_round rounds. This will reduce excessive iterations.
- **lambda** : lambda specifies regularization. Typical value ranges from 0 to 1.

- **min_gain_to_split** : This parameter will describe the minimum gain to make a split. It can be used to control the number of useful splits in the tree.
- **max_cat_group** : When the number of categories is large, finding the split point on it is easily over-fitting. So LightGBM merges them into 'max_cat_group' groups, and finds the split points on the group boundaries, default:64.

5. LightGBM implementation in Python

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import os
for dirname, _, filenames in os.walk(r"C:\Users\Hanshu\Desktop\kaggle_dataset"):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

C:\Users\Hanshu\Desktop\kaggle_dataset\Breast_cancer_data.csv
C:\Users\Hanshu\Desktop\kaggle_dataset\lightgbm-classifier-in-python.ipynb
C:\Users\Hanshu\Desktop\kaggle_dataset\naive-bayes-classifier-in-python.ipynb
C:\Users\Hanshu\Desktop\kaggle_dataset\adult.csv\adult.csv

```
In [4]: # ignore warnings
import warnings
warnings.filterwarnings('ignore')
```

Read dataset

```
In [5]: # Load and preview data

df = pd.read_csv(r'c:\Users\Hanshu\Desktop\kaggle_dataset\Breast_cancer_data.csv')
df.head()
```

```
Out[5]:
```

	mean_radius	mean_texture	mean_perimeter	mean_area	mean_smoothness	diagnos
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

View summary of dataset

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   mean_radius           569 non-null   float64
1   mean_texture          569 non-null   float64
2   mean_perimeter        569 non-null   float64
3   mean_area             569 non-null   float64
4   mean_smoothness       569 non-null   float64
5   diagnosis             569 non-null   int64
dtypes: float64(5), int64(1)
memory usage: 26.8 KB
```

- We can see that there are 6 columns in the dataset and there are no missing values.

Check the distribution of target variable

- target variable is `diagnosis`
- check the distribution of the target variable.

```
In [7]: # check the distribution of the target variable

df['diagnosis'].value_counts()
```

```
Out[7]: diagnosis
1      357
0      212
Name: count, dtype: int64
```

- The target variable is `diagnosis`. It contains 2 values - 0 and 1.
- `0` is for **Negative prediction** and `1` for **Positive prediction**.
- We can see that the problem is binary classification task.

Declare feature vector and target variable

```
In [9]: x = df[['mean_radius', 'mean_texture', 'mean_perimeter', 'mean_area', 'mean_smoothness']]
y = df['diagnosis']
```

Split dataset into training and test set

```
In [10]: # split the dataset into the training set and test set
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3, random_state = 42)
```

LightGBM Model Development and Training

- We need to convert our training data into LightGBM dataset format(this is mandatory for LightGBM training).
- After creating the necessary dataset, we created a python dictionary with parameters and their values.
- Accuracy of the model depends on the values we provide to the parameters.
- In the end block of code, we simply trained model with 100 iterations.

```
In [11]: # build the lightgbm model
import lightgbm as lgb
clf = lgb.LGBMClassifier()
clf.fit(x_train, y_train)
```

[illegible]


```
clf.fit(x_train, y_train)
```

```
#clf.get_params()
```

Out[20]:

```
LGBMClassifier
LGBMClassifier(n_estimator=100, verbose=-1)
```

In []:

```
from lightgbm import LGBMClassifier

#create a classifier
#clf = LGBMClassifier(n_estimator=100, boosting_type='gbdt', learning_rate=0.1,

#fit the model
clf.fit(x_train, y_train)

# print the model then it show parameters
#print(clf)
clf.get_params()
```

Out[]:

```
{'boosting_type': 'gbdt',
 'class_weight': None,
 'colsample_bytree': 1.0,
 'importance_type': 'split',
 'learning_rate': 0.1,
 'max_depth': -1,
 'min_child_samples': 20,
 'min_child_weight': 0.001,
 'min_split_gain': 0.0,
 'n_estimators': 100,
 'n_jobs': None,
 'num_leaves': 31,
 'objective': None,
 'random_state': 42,
 'reg_alpha': 0.0,
 'reg_lambda': 0.0,
 'subsample': 1.0,
 'subsample_for_bin': 200000,
 'subsample_freq': 0,
 'n_estimator': 100}
```

Model Prediction

In [19]:

```
# predict the results
y_pred = clf.predict(x_test)
```

View Accuracy

In [21]:

```
# view accuracy

from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_pred, y_test)
print('LightGBM Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_
```

LightGBM Model accuracy score: 0.9298

Here, `y_test` are the true class labels and `y_pred` are the predicted class labels in the test-set.

Compare train and test set accuracy

- Now, I will compare the train-set and test-set accuracy to check for overfitting.

```
In [ ]: # print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train,
```

```
In [26]: from sklearn.metrics import accuracy_score

# predictions
y_train_pred = clf.predict(x_train)
y_test_pred = clf.predict(x_test)

#accuracy
train_acc = accuracy_score(y_train, y_train_pred)
test_acc = accuracy_score(y_test, y_test_pred)

print("Training accuracy:", train_acc)
print('Testing accuracy:', test_acc)
```

Training accuracy: 1.0

Testing accuracy: 0.9298245614035088

Check for Overfitting

```
In [28]: # print the scores on training and test set

print('Training set score: {:.4f}'.format(clf.score(x_train, y_train)))

print('Test set score: {:.4f}'.format(clf.score(x_test, y_test)))
```

Training set score: 1.0000

Test set score: 0.9298

- The training and test set accuracy are quite comparable. So, we cannot say there is overfitting.

Confusion-matrix

```
In [29]: # view confusion-matrix
# Print the Confusion Matrix and slice it into four pieces

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print('Confusion matrix\n\n', cm)
print('\nTrue Positives(TP) = ', cm[0,0])
print('\nTrue Negatives(TN) = ', cm[1,1])
print('\nFalse Positives(FP) = ', cm[0,1])
print('\nFalse Negatives(FN) = ', cm[1,0])
```

Confusion matrix

```
[[ 55   8]
 [  4 104]]
```

True Positives(TP) = 55

True Negatives(TN) = 104

False Positives(FP) = 8


False Negatives(FN) = 4

```
In [30]: # visualize confusion matrix with seaborn heatmap

cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                          index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

Out[30]: <Axes: >



```
In [31]: # **Classification Metrics**
```

```
In [32]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.93	0.87	0.90	63
1	0.93	0.96	0.95	108
accuracy			0.93	171
macro avg	0.93	0.92	0.92	171
weighted avg	0.93	0.93	0.93	171

6. LightGBM Parameter Tuning

Table of Contents

- In this section, I will discuss some tips to improve LightGBM model efficiency.
- Following set of practices can be used to improve your model efficiency.
 - 1 **num_leaves** : This is the main parameter to control the complexity of the tree model. Ideally, the value of num_leaves should be less than or equal to $2^{(\text{max_depth})}$. Value more than this will result in overfitting.
 - 2 **min_data_in_leaf** : Setting it to a large value can avoid growing too deep a tree, but may cause under-fitting. In practice, setting it to hundreds or thousands is enough for a large dataset.
 - 3 **max_depth** : We also can use max_depth to limit the tree depth explicitly.

For Faster Speed

- Use bagging by setting `bagging_fraction` and `bagging_freq`.
- Use feature sub-sampling by setting `feature_fraction`.
- Use small `max_bin`.
- Use `save_binary` to speed up data loading in future learning.

For better accuracy

- Use large `max_bin` (may be slower).
- Use small `learning_rate` with large `num_iterations`
- Use large `num_leaves` (may cause over-fitting)
- Use bigger training data
- Try `dart`
- Try to use categorical feature directly.

To deal with over-fitting

- Use small `max_bin`

- Use small `num_leaves`
- Use `min_data_in_leaf` and `min_sum_hessian_in_leaf`
- Use bagging by set `bagging_fraction` and `bagging_freq`
- Use feature sub-sampling by set `feature_fraction`
- Use bigger training data
- Try `lambda_l1`, `lambda_l2` and `min_gain_to_split` to regularization
- Try `max_depth` to avoid growing deep tree